

LongAudioQA: Benchmarking Long-form Audio Meeting QA

Three Sub-corpora

Dataset	Dur.	#Q&A
AliMeeting	14.9 h	3,013
AMI Meeting	18.2 h	3,243
DailyTalk	21.6 h	10,200

Five Question Types

Factual

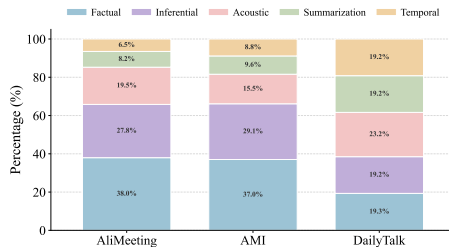
Inferential

Temporal

Summarization

Acoustic-Aware

Quality control: hybrid LLM + human annotation;
 $\kappa=0.91$; timestamp IoU > 0.9.



Distribution across all three corpora.

Unique contribution: **Acoustic-Aware** questions require grounding textual semantics with paralinguistic signals.